

# Introduction to Data Linkage

ID 529: Data Management and Analytic Workflows in R

Dean Marengi | Friday, January 16<sup>th</sup>, 2026

<https://id529.github.io/>

ID  
529

# Motivation

- We've discussed:
  - Data manipulation using `dplyr`
  - Leveraging `tidyverse` packages to work with specific types of data
    - Numeric, factors, text strings, dates/times, etc.
- However, we've only considered data manipulation in the context of individual datasets
  - We are often interested in combining data from multiple different files or sources
  - Need tools that help us work with and combine multiple datasets

<https://id529.github.io/>

# Learning objectives

- Understand the basic principles of data linkage
  - Common methods implemented for joining datasets
  - How methods differ, and why these differences should inform your approach
- Learn about different R functions available for data linkage
  - `dplyr` two-table verbs
- Learn how to implement `dplyr` two-table verbs to link datasets
  - Mutating joins
  - Filtering joins
  - Set operations

<https://id529.github.io/>

# Background

<https://id529.github.io/>

# Why do we need data linkage?

- It's rare that all the data we need for an analysis are contained within a single file
- Often interested in combining data across two or more files or sources
- For example, NHANES generates a number of datasets for each survey cycle
  - Separate files for demographics, clinical exams, laboratory tests, questionnaires, etc.
  - Participant data from a given survey cycle must be linked across multiple datasets with a unique identifier
  - **This is how the `ID529data::nhanes_id529` dataset was created!**

```
# A tibble: 14 × 2
  Data.File.Name Data.File.Description
  <chr>          <chr>
  1 BPX_J         Blood Pressure
  2 BMX_J         Body Measures
  3 OHXDEN_J     Oral Health - Dentition
  4 OHXREF_J     Oral Health - Recommendation of...
  5 DXXFEM_J     Dual-Energy X-ray Absorptiometr...
  6 DXX_J         Dual-Energy X-ray Absorptiometr...
  7 DXXSPN_J     Dual-Energy X-ray Absorptiometr...
  8 LUX_J         Liver Ultrasound Transient Elas...
  9 DXXAG_J      Dual-Energy X-ray Absorptiometr...
 10 BPXO_J       Blood Pressure - Oscillometric ...
 11 AUX_J        Audiometry
 12 AUXAR_J     Audiometry - Acoustic Reflex
 13 AUXTYM_J    Audiometry - Tympanometry
 14 AUXWBR_J    Audiometry - Wideband Reflectan...
```

```
# A tibble: 10 × 2
  Data.File.Name Data.File.Description
  <chr>          <chr>
  1 DR1TOT_J     Dietary Interview - Total Nutri...
  2 DR2TOT_J     Dietary Interview - Total Nutri...
  3 DR1IFF_J     Dietary Interview - Individual ...
  4 DR2IFF_J     Dietary Interview - Individual ...
  5 DS1IDS_J     Dietary Supplement Use 24-Hour ...
  6 DSQTOT_J     Dietary Supplement Use 30-Day ...
  7 DS2IDS_J     Dietary Supplement Use 24-Hour ...
  8 DS1TOT_J     Dietary Supplement Use 24-Hour ...
  9 DS2TOT_J     Dietary Supplement Use 24-Hour ...
 10 DSQIDS_J    Dietary Supplement Use 30-Day ...
```

# How do we link datasets together in R?

- There are many R functions that allow us to link data
  - `base` R includes the `merge()`, `rbind()`, `cbind()` functions
  - `dplyr` has many functions (verbs) for performing two-table operations
- We will focus on `dplyr`
  - `dplyr` functions tend to be more consistent, and intuitive to implement
  - Can be combined into chained data manipulation sequences (`|>` or `%>%`)
  - Also offers additional types of joins such as semi-joins and anti-joins

# dplyr two table verbs

Three classes of **dplyr** verbs that work with two tables at a time

- Mutating joins
  - Combine variables from multiple tables
  - Adds new columns to one table with matched rows from another table
  - Two types:
    - Inner join: `inner_join()`
    - Outer joins: `left_join()`, `right_join()`, `full_join()`
- Filtering joins
  - Filter rows from one table if they match a row in a second table
    - Unmatched rows are discarded if join condition is not met
  - Great tools for identifying row mismatches between tables
  - `semi_join()`, `anti_join`

<https://id529.github.io/>

<https://dplyr.tidyverse.org/articles/two-table.html>



# dplyr two table verbs (cont.)

## Set operations

- Combine observations from two datasets, but treats observations as set elements
  - i.e., treat rows in each table as individual items in a set
    - Both tables need to contain the same columns
    - Find rows in both tables (**intersection**)
    - Find rows in either table, but not both (**union**)
    - Find rows in one table, but not other (**difference**)
- `intersect()`, `union()`, `setdiff()`

<https://id529.github.io/>



# Example dataset

## Overview

- **NHANES** dataset available on the ID529 GitHub
- The dataset includes individual-level:
  - Demographic and clinical characteristics
  - Socioeconomic parameters
  - Blood measures of PFAS/PFOA
  - Dietary intake parameters
- For our examples, we will split these data into smaller datasets that each contain different variables, identifiers, and number of observations.
  - **clinical:** `id1, age, race_eth, sbp, ht, wt`
  - **pfas:** `id1, pfos, pfoa, pfna, pfhs, pfde`

```

1 # Create demog/clinical characteristics table (id1)
2 clinical <- nhanes %>%
3   rename(race_eth = race_ethnicity,
4         sbp = mean_BP,
5         ht = height,
6         wt = weight) %>%
7   select(id1, age, race_eth, sbp:ht)
8
9 # Create pfas data table (id1)
10 pfas <- nhanes %>%
11   select(id1, PFOS:PFDE) %>%
12   rename_with(~ str_to_lower(.)) %>%
13   filter(rowSums(is.na(.)) < 5)
```

<https://github.com/ID529/ID529data>

<https://id529.github.io/>

# Example dataset

```
1 # Demographic and clinical characteristics
2 clinical
```

```
# A tibble: 2,339 × 6
  id1     age race_eth          sbp    wt    ht
  <chr> <int> <fct>        <dbl> <dbl> <dbl>
1 73568     26 Non-Hispanic White 105.  47.1 152.
2 73571     76 Non-Hispanic White 126   102. 172.
3 73574     33 <NA>            121.  56.8 158
4 73576     16 Non-Hispanic Black 109.  67.3 170.
5 73577     32 Hispanic         119.  79.7 166.
6 73578     18 Hispanic         123.  109. 175.
7 73584     13 Non-Hispanic White 109.  53.1 145.
8 73587     14 <NA>            112   110. 169.
9 73597     50 Non-Hispanic Black NaN   104. 180.
10 73598    20 Hispanic         112   86.7 165
# i 2,329 more rows
```

```
1 # PFAS data filtered for obs with data
2 pfas
```

```
# A tibble: 2,168 × 6
  id1     pfos    pfoa    pfna    pfhs    pfde
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 73568     2.2     3     0.5     3     0.2
2 73571    10.2    4.77    1.3     2     0.3
3 73574     NA      NA     0.7     0.2     0.1
4 73576     4.7    2.37    0.6     7.6     0.2
5 73577     3     1.47    0.4     1.2     0.1
6 73584     7     2.37    0.8     0.8     0.2
7 73587    35.5    6.17    3.3     6.3     1.7
8 73598     4.7    1.8     0.5     1.6     0.2
9 73599     4.5    1.87    1.7     0.8     0.2
10 73600    6.3    1.67    0.5     1.6     0.2
# i 2,158 more rows
```

<https://id529.github.io/>

# dplyr mutating joins

<https://id529.github.io/>

# dplyr::left\_join()

## Main arguments

```

1 left_join(
2   x, # dataset 1 (i.e., the dataset on the left side)
3   y, # dataset 2 (i.e., the dataset on the right side)
4   by = NULL,
5   suffix = c(".x", ".y"),
6   keep = FALSE
7 )


```

- **x** = Data frame 1 (left table)
- **y** = Data frame 2 (right table)
- **by** = Column names(s) of variables to match rows by (key)
- **suffix** = Character string appended to column names that appear in both **x** and **y**
- **keep**= If **TRUE**, preserve the join keys from both **x** and **y**

Gif Source: <https://github.com/gadenbuie/tidyexplain>

## Description

- Adds columns from **y** (right table) to **x** (left table)
- All rows from **x** are preserved in the join
  - i.e., keeps all rows from **x** regardless of whether there's a match in **y**
- Of the joins, used most frequently
  - Allow you to add new variables from other tables



# dplyr::left\_join() example

```
1 # Print clinical dataset
2 clinical
```

```
# A tibble: 2,339 × 6
  id1    age race_eth      sbp     wt     ht
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispanic White  105.   47.1   152.
2 73571    76 Non-Hispanic White  126    102.   172.
3 73574    33 <NA>            121.   56.8   158
4 73576    16 Non-Hispanic Black 109.   67.3   170.
5 73577    32 Hispanic         119.   79.7   166.
6 73578    18 Hispanic         123.   109.   175.
7 73584    13 Non-Hispanic White 109.   53.1   145.
8 73587    14 <NA>            112    110.   169.
9 73597    50 Non-Hispanic Black NaN    104.   180.
10 73598   20 Hispanic         112    86.7   165
# i 2,329 more rows
```

```
1 # Print PFAS dataset
2 pfas
```

```
# A tibble: 2,168 × 6
  id1    pfos   pfoa   pfna   pfhs   pfde
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 73568    2.2    3     0.5    3     0.2
2 73571   10.2   4.77   1.3    2     0.3
3 73574    NA     NA    0.7    0.2    0.1
4 73576    4.7   2.37   0.6    7.6    0.2
5 73577    3     1.47   0.4    1.2    0.1
6 73584    7     2.37   0.8    0.8    0.2
7 73587   35.5   6.17   3.3    6.3    1.7
8 73598    4.7   1.8    0.5    1.6    0.2
9 73599    4.5   1.87   1.7    0.8    0.2
10 73600   6.3   1.67   0.5    1.6    0.2
# i 2,158 more rows
```

```
1 # Perform a left join
2 left_join(x = clinical, y = pfas, by = "id1")
```

```
# A tibble: 2,339 × 11
  id1    age race_eth      sbp     wt     ht   pfos
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl> <dbl>
1 73568    26 Non-Hispan...  105.   47.1   152.   2.2
2 73571    76 Non-Hispan...  126    102.   172.   10.2
3 73574    33 <NA>          121.   56.8   158    NA
4 73576    16 Non-Hispan...  109.   67.3   170.   4.7
5 73577    32 Hispanic       119.   79.7   166.   3
6 73578    18 Hispanic       123.   109.   175.   NA
7 73584    13 Non-Hispan...  109.   53.1   145.   7
8 73587    14 <NA>          112    110.   169.   35.5
9 73597    50 Non-Hispan...  NaN    104.   180.   NA
10 73598   20 Hispanic       112    86.7   165    4.7
# i 2,329 more rows
# i 4 more variables: pfoa <dbl>, pfna <dbl>,
#   pfhs <dbl>, pfde <dbl>
```

# dplyr::right\_join()

## Main arguments

```

1 right_join(
2   x,
3   y,
4   by = NULL,
5   suffix = c(".x", ".y"),
6   keep = FALSE
7 )


```

- **x** = Data frame 1 (left table)
- **y** = Data frame 2 (right table)
- **by** = Column names(s) of variables to match rows by (key)
- **suffix** = Character string appended to column names that appear in both **x** and **y**
- **keep**= If **TRUE**, preserve the join keys from both **x** and **y**

Gif Source: <https://github.com/gadenbuie/tidyexplain>

## Description

- Adds columns from **x** (left table) to **y** (right table)
- All rows from **y** are preserved in the join
  - i.e., keeps all rows from **y** regardless of whether there's a match in **x**



# dplyr::right\_join() example

```
1 # Print clinical dataset
2 clinical
```

```
# A tibble: 2,339 × 6
  id1    age race_eth      sbp     wt     ht
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispanic White  105.   47.1   152.
2 73571    76 Non-Hispanic White  126    102.   172.
3 73574    33 <NA>            121.   56.8   158
4 73576    16 Non-Hispanic Black 109.   67.3   170.
5 73577    32 Hispanic         119.   79.7   166.
6 73578    18 Hispanic         123.   109.   175.
7 73584    13 Non-Hispanic White 109.   53.1   145.
8 73587    14 <NA>            112    110.   169.
9 73597    50 Non-Hispanic Black NaN    104.   180.
10 73598   20 Hispanic         112    86.7   165
# i 2,329 more rows
```

```
1 # Print PFAS dataset
2 pfas
```

```
# A tibble: 2,168 × 6
  id1    pfos    pfoa    pfna    pfhs    pfde
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 73568    2.2     3     0.5     3     0.2
2 73571   10.2    4.77    1.3     2     0.3
3 73574    NA      NA     0.7     0.2     0.1
4 73576    4.7    2.37    0.6     7.6     0.2
5 73577    3     1.47    0.4     1.2     0.1
6 73584    7     2.37    0.8     0.8     0.2
7 73587   35.5   6.17    3.3     6.3     1.7
8 73598    4.7    1.8     0.5     1.6     0.2
9 73599    4.5    1.87    1.7     0.8     0.2
10 73600   6.3    1.67    0.5     1.6     0.2
# i 2,158 more rows
```

```
1 right_join(x = clinical, y = pfas, by = "id1")
```

```
# A tibble: 2,168 × 11
  id1    age race_eth      sbp     wt     ht    pfos    pfoa    pfna    pfhs    pfde
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispan...  105.   47.1   152.   2.2
2 73571    76 Non-Hispan...  126    102.   172.   10.2
3 73574    33 <NA>          121.   56.8   158    NA
4 73576    16 Non-Hispan...  109.   67.3   170.   4.7
5 73577    32 Hispanic       119.   79.7   166.   3
6 73584    13 Non-Hispan...  109.   53.1   145.   7
7 73587    14 <NA>          112    110.   169.   35.5
8 73598    20 Hispanic       112    86.7   165    4.7
9 73599    13 Non-Hispan...  118    44.9   159.   4.5
10 73600   37 Non-Hispan...  149.   126.   185.   6.3
# i 2,158 more rows
# i 4 more variables: pfoa <dbl>, pfna <dbl>,
#   pfhs <dbl>, pfde <dbl>
```

# dplyr::inner\_join()

## Main arguments

```

1 inner_join(
2   x,
3   y,
4   by = NULL,
5   copy = FALSE,
6   suffix = c(".x", ".y"),
7   ...,
8   keep = FALSE
9 )

```

- **x** = Data frame **x**
- **y** = Data frame **y**
- **by**= Column name(s) to join **x** and **y** by
- **suffix**= Suffix to append to duplicate vars
- **keep**= Keep the join key from both **x** and **y**

## Description

- Returns rows that match in **both x and y**
  - Unmatched rows are dropped
- Most restrictive of the mutating joins

inner\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

<https://id529.github.io/>

Gif Source: <https://github.com/gadenbuie/tidyexplain>

# dplyr::inner\_join() example

```
1 # Print clinical dataset
2 clinical
```

```
# A tibble: 2,339 × 6
  id1    age race_eth      sbp     wt     ht
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispanic White  105.   47.1   152.
2 73571    76 Non-Hispanic White  126    102.   172.
3 73574    33 <NA>            121.   56.8   158
4 73576    16 Non-Hispanic Black 109.   67.3   170.
5 73577    32 Hispanic         119.   79.7   166.
6 73578    18 Hispanic         123.   109.   175.
7 73584    13 Non-Hispanic White 109.   53.1   145.
8 73587    14 <NA>            112    110.   169.
9 73597    50 Non-Hispanic Black NaN    104.   180.
10 73598   20 Hispanic         112    86.7   165
# i 2,329 more rows
```

```
1 # Print PFAS dataset
2 pfas
```

```
# A tibble: 2,168 × 6
  id1    pfos pfoa pfna pfhs pfde
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 73568    2.2  3     0.5   3     0.2
2 73571   10.2  4.77  1.3   2     0.3
3 73574    NA    NA    0.7   0.2   0.1
4 73576    4.7   2.37  0.6   7.6   0.2
5 73577    3     1.47  0.4   1.2   0.1
6 73584    7     2.37  0.8   0.8   0.2
7 73587   35.5  6.17  3.3   6.3   1.7
8 73598    4.7   1.8   0.5   1.6   0.2
9 73599    4.5   1.87  1.7   0.8   0.2
10 73600   6.3   1.67  0.5   1.6   0.2
# i 2,158 more rows
```

```
1 # Perform an inner join
2 inner_join(x = clinical, y = pfas, by = "id1")
```

```
# A tibble: 2,168 × 11
  id1    age race_eth      sbp     wt     ht   pfos
  <chr> <int> <fct>     <dbl>   <dbl>   <dbl> <dbl>
1 73568    26 Non-Hispan...  105.   47.1   152.   2.2
2 73571    76 Non-Hispan...  126    102.   172.   10.2
3 73574    33 <NA>          121.   56.8   158    NA
4 73576    16 Non-Hispan...  109.   67.3   170.   4.7
5 73577    32 Hispanic     119.   79.7   166.   3
6 73584    13 Non-Hispan...  109.   53.1   145.   7
7 73587    14 <NA>          112    110.   169.   35.5
8 73598    20 Hispanic     112    86.7   165    4.7
9 73599    13 Non-Hispan...  118    44.9   159.   4.5
10 73600   37 Non-Hispan... 149.   126.   185.   6.3
# i 2,158 more rows
# i 4 more variables: pfoa <dbl>, pfna <dbl>,
#   pfhs <dbl>, pfde <dbl>
```

# dplyr::full\_join()

## Main arguments

```

1 full_join(
2   x,
3   y,
4   by = NULL,
5   suffix = c(".x", ".y"),
6   keep = FALSE
7 )

```

- **x** = Data frame **x**
- **y** = Data frame **y**
- **by**= Column name(s) to join **x** and **y** by
- **suffix**= Suffix to append to duplicate vars
- **keep**= Keep the join key from both **x** and **y**

## Description

- Returns all rows in **x** or **y**
- Least restrictive
  - No observations are removed

full\_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

<https://id529.github.io/>

Gif Source: <https://github.com/gadenbuie/tidyexplain>

# dplyr::full\_join() example

```
1 # Print clinical dataset
2 clinical
```

```
# A tibble: 2,339 × 6
  id1    age race_eth      sbp     wt     ht
  <chr> <int> <fct>    <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispanic White  105.   47.1   152.
2 73571    76 Non-Hispanic White  126    102.   172.
3 73574    33 <NA>            121.   56.8   158
4 73576    16 Non-Hispanic Black 109.   67.3   170.
5 73577    32 Hispanic         119.   79.7   166.
6 73578    18 Hispanic         123.   109.   175.
7 73584    13 Non-Hispanic White 109.   53.1   145.
8 73587    14 <NA>            112    110.   169.
9 73597    50 Non-Hispanic Black NaN    104.   180.
10 73598   20 Hispanic         112    86.7   165
# i 2,329 more rows
```

```
1 # Print PFAS dataset
2 pfas
```

```
# A tibble: 2,168 × 6
  id1    pfos    pfoa    pfna    pfhs    pfde
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 73568    2.2     3     0.5     3     0.2
2 73571   10.2    4.77    1.3     2     0.3
3 73574    NA      NA     0.7     0.2     0.1
4 73576    4.7    2.37    0.6     7.6     0.2
5 73577    3     1.47    0.4     1.2     0.1
6 73584    7     2.37    0.8     0.8     0.2
7 73587   35.5   6.17    3.3     6.3     1.7
8 73598    4.7    1.8     0.5     1.6     0.2
9 73599    4.5    1.87    1.7     0.8     0.2
10 73600   6.3    1.67    0.5     1.6     0.2
# i 2,158 more rows
```

```
1 # Perform full join
2 full_join(x = clinical, y = pfas, by = "id1")
```

```
# A tibble: 2,339 × 11
  id1    age race_eth      sbp     wt     ht    pfos    pfoa    pfna    pfhs    pfde
  <chr> <int> <fct>    <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
1 73568    26 Non-Hispan...  105.   47.1   152.   2.2
2 73571    76 Non-Hispan...  126    102.   172.   10.2
3 73574    33 <NA>            121.   56.8   158    NA
4 73576    16 Non-Hispan...  109.   67.3   170.   4.7
5 73577    32 Hispanic       119.   79.7   166.    3
6 73578    18 Hispanic       123.   109.   175.   NA
7 73584    13 Non-Hispan...  109.   53.1   145.    7
8 73587    14 <NA>            112    110.   169.   35.5
9 73597    50 Non-Hispan...  NaN    104.   180.   NA
10 73598   20 Hispanic       112    86.7   165    4.7
# i 2,329 more rows
# i 4 more variables: pfoa <dbl>, pfna <dbl>,
#   pfhs <dbl>, pfde <dbl>
```

# dplyr filtering joins

<https://id529.github.io/>

# dplyr::semi\_join()


## Main arguments

```
1 semi_join(x,
2           y,
3           by = NULL,
4 )
```

- **x** = Data frame 1 (left)
- **y** = Data frame 2 (right)
- **by** = Column names(s) of variables to match rows by (key)

## Description

- Keeps all observations in **x** that have a match in **y**
- Useful for filtering **x** by the presence of a match in **y**



Gif Source: <https://github.com/gadenbuie/tidyexplain>

# dplyr::anti\_join()


## Main arguments

```
1 anti_join(x,
2           y
3           by = NULL
4 )
```

- **x** = Data frame 1 (left)
- **y** = Data frame 2 (right)
- **by** = Column names(s) of variables to match rows by (key)

## Description

- Returns rows from **x** where there is no match in **y**
- Useful when you want to filter the first table by the absence of a match in the second table
  - That is, identify observations in **x** that don't appear in **y**



Gif Source: <https://github.com/gadenbuie/tidyexplain>

# dplyr set operations

<https://id529.github.io/>

# dplyr::union()

## Main arguments

```
1 union(x, y)
2 union_all(x, y)
```

- **x** = Data frame 1 (left)
- **y** = Data frame 2 (right)

## Description

- **union()**
  - Return unique observations in **x** and **y**
- **union\_all()**
  - Return unique and duplicate observations in **x** and **y**

union(x, y)	
1	a
1	b
2	a

1	a
2	b

union_all(x, y)	
1	a
1	b
2	a

1	a
2	b

<https://id529.github.io/>

Gif Source: <https://github.com/gadenbuie/tidyexplain>

# dplyr::setdiff()


## Main arguments

```
1 setdiff(x, y, ...)
```

- **x** = Data frame 1 (left)
- **y** = Data frame 2 (right)

## Description

- Returns the rows that are in the first table but not in the second
  - That is, return observations in x, but not in y



Gif Source: <https://github.com/gadenbuie/tidyexplain>

# Key takeaways

- Joining together two or more datasets is a routine component of data analysis workflows
  - Different types of joins are useful in different situations
  - It's important to be aware of how and when to use each (left, right, inner, etc.)
    - All have distinct implications for how the underlying data are manipulated
- Like all data manipulation tasks, joining data can get complicated!
- Common challenges include:
  - Inconsistencies in unique identifiers
  - Duplicate records and missing values
  - Non-unique identifiers (require matching data frames on two or more variables)
  - Joining datasets can be complex and requires a good understanding of the data structure
- **dplyr** includes a number of two-table functions (verbs)
  - Like other **dplyr** verbs, two-table verbs are:
    - Intuitive to use
    - Can be incorporated into a sequence of data cleaning operations using the pipe operator

<https://id529.github.io/>

# Resources

- **dplyr Two-table verbs vignette:** <https://dplyr.tidyverse.org/articles/two-table.html>
- **R for Data Science - Relational Data:** <https://r4ds.had.co.nz/relational-data.html>
- **The Epidemiologist R Handbook:** [https://epirhandbook.com/en/new\\_pages/joining\\_matching.html](https://epirhandbook.com/en/new_pages/joining_matching.html)

<https://id529.github.io/>

# Activity

- **Clone the repository here:** [https://github.com/dmarengi/joining\\_data](https://github.com/dmarengi/joining_data)
- **ID 529 dataset details:** <https://github.com/ID529/ID529data>

<https://id529.github.io/>